

Audit & Optimization of the Calm.Profile Diagnostic System

1. Streamlining the Diagnostic Report Structure

The current calm.profile report spans ~12 pages, covering everything from an executive summary to detailed findings and recommendations. To improve clarity and impact, we recommend restructuring it following a classic consulting format: **Cover → Executive Summary → Background/Context → Key Findings → Recommendations → Appendix** ¹. This ensures a logical flow that busy founders can easily follow. Start with a brief **Introduction/Background** section that frames the agency's challenges ("chaos vs. calm") and why the diagnostic was conducted. Leading consultancies emphasize stating the problem and context up front to hook the reader ² ³. In calm.profile, this could be a short paragraph summarizing the agency's current state (e.g. "frequent coordination fire-drills, low creative time, tool overload") and the assessment scope.

Next, strengthen the **Executive Summary** using a **BLUF (Bottom Line Up Front)** approach ⁴. Open with the *most important insight* – for example: "**Bottom Line:** Your agency operates as a **Balanced Architect** archetype, but is losing ~70 hours/month to operational friction, costing an estimated \$81K/year in capacity." Present 2-3 high-level findings and the top recommendations in bullet form ⁵. The current summary already lists key stats (overhead hours, coordination load, etc.), which is excellent. We suggest adding a concise narrative sentence tying those stats to an overall opportunity or risk (e.g. "These inefficiencies equate to a ~37% productivity boost opportunity if addressed"). This delivers a one-page story of the diagnosis, increasing perceived value for a time-crunched founder ⁶ ⁷. Make sure to keep the tone **consultative and founder-friendly** – use plain language and avoid jargon. In fact, syris's brand positioning notes that the messaging should be straightforward and speak directly to founders ⁸. Maintaining that neutral, advisory tone builds credibility ("diagnose before you prescribe"), aligning with best practices.

Within the main report, consider consolidating and simplifying where possible. For example, the template currently has separate sections for "Operational Inefficiency Findings" and a "Friction Deep Dive" by domain. To streamline, you might **merge related findings and domains** so that each major finding is immediately followed by deeper analysis of its root causes in the friction domains. This way, the report flows as a clear story: finding → evidence → root cause → related domain detail → solution. Ensure each major finding explicitly ties to a proposed solution or improvement (consulting reports should "**introduce solutions to the client's specific problems**", not just diagnostics ⁹ ¹⁰). You can still keep the structured breakdown (Communication, Task Handoff, etc.), but perhaps move any non-critical domain details to an appendix or sidebars. This keeps the core narrative focused on the biggest pain points.

It's also important to **highlight key facts visually or call them out** so they stand out. Many consulting deliverables use call-out boxes or bold infographics for impactful stats ¹¹ ⁵. In a 12-page report, you might introduce a one-page "**Key Metrics Snapshot**" after the exec summary – for instance, a small table or infographic summarizing "Weekly hours lost per person", "% time in non-billable work vs creative work",

“Potential \ \$ savings” etc. Presenting a few quantified pain points in a standalone section helps busy readers quickly grasp the situation ⁶ ⁷ . The existing report already quantifies lost hours and dollars; surfacing those in a visually distinct way (e.g. an infographic with icons for time and money) will enhance perceived value.

Lastly, apply a **coherent story thread throughout the report**. Use the archetype as a red thread: for example, “As a **‘Project Architect’** style leader (Workstyle Profile section), you excel at structure but may over-index on control – this is reflected in Finding #1, where rigid task workflows cause delays.” This connects the dots for the client and reinforces that the diagnostic isn’t just random data points but a narrative about their business. Every section should answer “so what?” for the founder. If any page or metric doesn’t clearly ladder up to an insight or recommendation, consider trimming it. A tight storyline improves clarity and keeps the founder engaged. In summary, by reordering sections for logic, highlighting key insights up front, and ensuring each finding ties to actionable advice, the report will feel more impactful and easier to digest – much like a top-tier consultancy deliverable rather than a generic audit.

2. Enhancing the Scoring Model (Traits-to-Profile Algorithm)

The calm.profile assessment uses 20 questions to map the team’s workstyle across four axes – Structure Orientation, Collaboration Mode, Scope Focus, and Tempo Preference – ultimately generating an “operational archetype” profile. To ensure this scoring model is reliable, we recommend a two-pronged review: **statistical reliability testing** and potential simplification.

Reliability & Consistency: Verify that each of the four axes is being measured consistently by its set of questions. With ~5 questions per axis, you can calculate an internal consistency metric (e.g. **Cronbach’s alpha**) for each trait scale to see if those questions indeed correlate well (commonly, $\alpha > 0.7$ is desirable for a formative assessment). If an axis’s questions show low consistency, it may indicate some questions are off-target or interpreted inconsistently by respondents. You might refine or replace those items with clearer ones. For example, if “Scope Focus” includes questions that mix unrelated concepts, respondents’ answers could vary too widely. Each question should clearly target one dimension at a time (e.g. *“I prefer to dive into project details rather than delegate big-picture planning”* strictly addresses macro vs. micro focus). Ensuring **each trait’s questions are unidimensional** and easy to interpret will improve the confidence in the scores ¹² ¹³ . Additionally, consider conducting a short pilot: have a few users take the survey twice (test-retest) or have multiple team members respond and see if the archetype results make intuitive sense. If the “confidence band” for the type (as mentioned in the code) is often coming out “low”, that’s a signal to tighten the questions or scoring thresholds for more distinct outcomes.

Simplification & Accuracy: The current algorithm likely sums or averages responses to get axis scores (0–100 scale) and then determines an archetype from the combination. If this trait-to-profile mapping is overly complex (for example, lots of conditional logic in the JavaScript), see if it can be **simplified without losing meaning**. Often a simpler model (even a straight scoring or a 2x2 matrix) is more transparent to clients and easier to maintain. For instance, if each axis is essentially a continuum between two styles (structured vs flexible, synchronous vs async, etc.), you could define the archetype as simply the highest/most pronounced dimension or a combination of top two dimensions. A complex ruleset (like “if Structure >70 and Tempo <50 then Type A, else if...”) might produce nuanced types but could be hard to validate. It may be better to have fewer archetypes that are **clearly distinct**. This also helps the founder identify with their profile readily. If needed, provide an option for a “hybrid” profile when scores are balanced – the code’s `confidence_band` logic hints at flagging low confidence results and even naming a hybrid of primary-secondary type ¹⁴ ¹⁵ .

This is a good practice to acknowledge nuance. Just ensure the criteria for hybrid vs primary type are well-defined and tested (e.g. within 5-point difference on an axis might trigger a hybrid label).

Automation of Scoring: Since the questions are in HTML/JS now, make sure the scoring calculations are fully automated and error-free. Ideally, the responses should feed directly into the Python backend or whatever system generates the report, rather than any manual transfer. If not already, use a script to map each survey answer to the appropriate axis score behind the scenes. This reduces the chance of human error in computing scores and enables you to quickly iterate on the model (you can tweak weightings or questions and instantly see updated results). Given the small number of questions, a simple weighting (all questions equal weight on their axis) is likely fine. But if you discover certain questions are stronger indicators, you could weight them slightly more – again, something easily adjusted in code if the pipeline is automated.

Calibration and Benchmarking: To improve the **founder's trust in the scores**, continue using industry benchmarks (the report template shows industry avg of 65 for Structure, 55 for Collaboration, etc.). These lend context and perceived validity. If possible, update these benchmarks over time with more data from agencies that take the assessment. Also consider adding brief text interpretations for extreme scores to give qualitative color (e.g. "Collaboration Mode 85/100 – very high, indicating a strong preference for synchronous, real-time teamwork, which can sometimes overwhelm those who prefer async communication"). This blends the raw score into a meaningful narrative and can be generated automatically based on thresholds.

In summary, the scoring model appears sound in design (covering four key dimensions of work style), but a reliability audit can validate its effectiveness. Simplifying the archetype logic to be transparent will also help both automation and client understanding. By automating calculations and refining question clarity, the assessment will yield consistent, credible profiles that form a rock-solid basis for the report's insights.

3. Increasing Automation in Report Generation

Generating the final PDF report should be as hands-free as possible, so you can scale calm.profile without drowning in manual work. Currently, the process involves Python templates and presumably some manual editing of the Markdown (or JSON) outputs to insert client-specific details. There are several opportunities to streamline this:

Use Templating for Dynamic Content: The provided Python snippet `DynamicReportGenerator` suggests you're already constructing a `report_content` data structure with all the numbers, narratives, and recommendations filled in ¹⁶ ¹⁷. Leverage this by binding it directly to the report template. For example, if you maintain the report in Markdown or HTML, use a templating engine like **Jinja2** to insert values (e.g., `{{ executive_summary.opening }}` for the opening line). This way, once the assessment data is scored, a script can automatically populate all the placeholders (agency name, scores, findings text, etc.) in one go. Medium's engineering blog gives a concise example of this pattern: define a template with placeholders and pair it with a JSON data source to output a filled report ¹⁸ ¹⁹. Your case is essentially the same; you have the data (scores, archetype, improvement stats) and a template (the 12-page structure). By automating the merge, you eliminate hand-copying values or writing repetitive text.

Automate Narrative Generation: Much of the narrative content can be pre-written and chosen based on rules – and indeed your code shows dictionaries for archetype-specific text and friction-point language. Continue extending this so that the system selects paragraphs or bullet points automatically. For instance, have a library of **Archetype Strengths/Blindspots** statements and pull in the ones for the identified archetype. The `ARCHETYPE_CONTENT` and `FRICTION_CONTENT` in your code already outline this approach, where each profile type and each friction domain has templated text for different score levels ²⁰ ²¹. This allows the PDF to be generated with minimal manual writing per client. You'll just need to sanity-check that the assembled text reads smoothly (maybe a light human edit or QA pass before sending, especially on early uses, but no heavy lifting). Over time, as you trust the automation, the entire narrative could be generated in one click.

Leverage Markdown/HTML to PDF Tools: Instead of manually tweaking a Word doc or InDesign file, adopt a scriptable conversion. Since your content is structured, you can use a library like **WeasyPrint** (HTML/CSS to PDF) or **Pandoc** (Markdown to PDF via LaTeX) to generate a polished PDF at the push of a button. For example, one team reported success using Jinja to create an HTML report and WeasyPrint to output it as a styled PDF, achieving professional layouts with much less effort ²² ²³. WeasyPrint + HTML/CSS is powerful because you can design the report like a web page (with your minimal calm.studio branding) and get a PDF that preserves those styles. This also makes it easier to include visuals (charts, icons) and consistent formatting. The Medium case study noted that using Jinja+WeasyPrint increased their report generation speed by ~75% and let them reuse template sections easily ²⁴ ²⁵. In your context, that means you could maintain, say, a “base template” for the report and include optional sections (perhaps skip a section if a certain condition isn't met – e.g., if no friction domain is critical, maybe the “emergency actions” box is omitted). A templating approach can handle such conditional content (indeed your code's `get_conditional_content` function is already determining flags like `show_emergency_actions` based on scores ²⁶ ²⁷). The automation should insert or hide those elements without you manually deleting things.

Minimize Manual Data Gathering: Automating the report also involves streamlining how input data is collected. If not already, host the 20-question survey online (even a Google Form or a simple web form) and have responses feed into your system directly. If you gather operational data (team size, tool list, etc.) via an intake form or interview, standardize those inputs as well. For example, you could have a small questionnaire or spreadsheet for the founder to fill in basic metrics (# of employees, avg meetings per week, etc.). Having those in a structured format (CSV or form submission) means you can parse them with Python and plug into the calculations (e.g., coordination load = meetings hours per week / total hours). Every piece of data that currently requires you to type it into the template is an opportunity for automation.

Testing and Iteration: Once the automated pipeline is set up (survey responses → scoring code → content generation → PDF formatting), test it on past cases. Make a “dummy” agency profile and run it end-to-end, then inspect the PDF to ensure everything lands in the right place (check for placeholder misses, line breaks, etc.). It's normal to fine-tune the template styling or phrasing after seeing the auto-generated result. But the payoff is big: the next real client's report might be produced in hours instead of days. You can then spend your time on the **analysis and recommendations** (the high-value consulting work) rather than on formatting. Plus, a highly automated process reduces errors and ensures consistency across reports. Clients will get the same quality and format every time, reinforcing syris's professional image.

In summary, **fully automate the report assembly** using your Python templating system, Jinja placeholders, and PDF generation tools. The goal is that once a client completes the diagnostic (and

perhaps a brief follow-up Q&A), you run a script and get a ready-to-send PDF that might only need minimal polishing. This will make calm.profile a scalable, repeatable product – crucial for founder appeal, since you can then confidently promise a quick turnaround (“a few days” as per your positioning) and potentially even lower cost or higher margin. Automation is key to delivering consistency and speed, which in turn reinforces the value to the client (they get insights fast) and frees you to focus on deeper advisory discussions instead of report assembly.

4. Upgrading Visuals and Data Presentation

Right now, the report contains some tables and text descriptions for scores and friction points. Enhancing the visual presentation will improve comprehension and give the report that “consultancy grade” polish. Founders will perceive a higher value when insights are illustrated clearly. We suggest the following upgrades:

Figure: Example radar chart comparing a sample agency's workstyle scores to industry benchmarks. Such visualizations make abstract concepts (structure, collaboration, scope, tempo) more tangible at a glance.

Operational Workstyle Charts: Instead of (or in addition to) a simple table of the four axis scores, present them in a **visual chart format**. A radar (spider) chart is a great choice for showing an agency's profile across the 4 axes relative to the industry average or an “ideal” benchmark. In the example above, the blue shape might represent the client and the dotted line the benchmark. It immediately shows where the client deviates – for instance, perhaps they score much higher on “Structure” but lower on “Tempo” than the norm. This gives a quick visual cue of the workstyle balance, which is more impactful than numbers alone. Management consultants often favor such charts to let the client self-identify gaps (“we’re way above average on structure – maybe we’re over-engineering things?”) without heavy text ²⁸. The radar chart can be generated automatically from the scores and embedded in the PDF. If radar charts are unfamiliar to some, an alternative is a simple bar graph: four bars for the client vs four bars for the benchmark on each axis. The key is to **use color or icons** to label each axis clearly (perhaps an icon for each: e.g. a gears icon for Structure, a chat icon for Collaboration, magnifying glass for Scope detail vs telescope for big-picture, a clock for Tempo). Visual cues help reinforce the concepts without the client needing to recall definitions from text.

Friction Heatmaps: The report template mentions a heat map of friction across stages (Lead Gen → Archive). Turning this into a true visual would greatly aid understanding. One idea is a **process timeline with color-coding**. You can draw the major workflow stages as a horizontal flow chart (or even just list them in a row) and use traffic-light colors under each to indicate friction intensity (green = low, yellow = medium, red = high). For example, if “Design” and “Review” are high-friction stages (as in the template example), those boxes or circles could be red, whereas “Lead Gen” might be green for low friction ²⁹. This heatmap instantly shows *where* in the process the biggest bottlenecks lie. It’s much easier for a client to digest than reading paragraphs about each stage. We suggest also adding a legend or labels (maybe percentage of inefficiency or a score out of 100 for each stage’s friction). A visual heatmap ties the analysis together spatially – the client can literally point to the diagram and see “ah, the creative review stage is our pain point.” It transforms abstract survey results into a concrete map of their workflow.

Status Dashboard for Domains: In the “Friction Domain Deep Dive” section, you currently list each domain with a score and status (Critical/Warning/Stable). Consider presenting this as a simple **dashboard graphic or table with color indicators**. For example, a table where each domain is a row, with columns for Score and Status, and a colored dot or symbol for status (● stable, ● caution, ● critical). Even without reading the

text, a founder can see perhaps two reds (e.g. **Decision Latency: 75/100 Critical**) and know immediately where to pay attention. You might also use small bar gauges next to each domain score – a bar partially filled to 75% with red coloring for a 75/100 score, etc. Visualizing these domain scores helps break up the text and gives a quick reference summary that complements the detailed bullet points you provide. It's essentially applying a **dashboard approach** within the report, which adds a professional touch (clients often love seeing “dashboard” elements, as it feels like an ongoing tool they can monitor). Remember, the goal is to *highlight the main data points graphically*, so the reader can grasp them even when skimming ⁷

30 .

ROI and Impact Graphics: The financial impact section is perfect for visual treatment. We suggest including the **5-Year Cumulative Impact** as a simple line or area chart. Plot the potential savings over time if improvements are made – year by year up to five years (as your template outlines). This could be a small line graph with the X-axis as Years 1–5 and Y-axis as dollars saved, perhaps with two lines: one for the “Realistic scenario” and one flat line showing “If no changes: \ \$0 saved” for contrast. Such a chart taps into the client's visual intuition about compounding gains (e.g. a steadily rising curve reaching, say, \ \$500K by Year 5). It drives home the value of acting now. Similarly, a pie chart or icon could illustrate the “ROI multiple” – e.g. a graphic showing \$495 investment vs ~\$50K savings, to emphasize the ROI (100x, etc.). While these numbers are already in text, visuals will make them *pop*. For quick wins vs long-term initiatives, you might use a **2x2 matrix** graphic for the impact/effort matrix mentioned. Plot some key recommendations as dots in the quadrants (Quick Wins = high impact/low effort, etc.), or at least present that matrix as a colored diagram. This not only looks professional (consultants love 2x2 matrices) but also helps the founder see why certain actions are prioritized first ¹⁰ .

Tool Consolidation Illustration: In the tool stack section, listing current vs recommended tools in a table is useful. To enhance it, you could add logos of the tools (if appropriate and not overly busy) or at least icons for categories (project management, communication, storage, etc.). Even a small before-and-after diagram could work: e.g., a cluttered cluster of 8 tool logos on the left labeled “Current: 8 tools” and a simplified cluster of 3 logos on the right labeled “Proposed: 3 tools”. This provides a **visual story of simplification**, which resonates with the promise of calm. Since your brand ethos is minimalism, showing a visual simplification (many to few) reinforces that. It gives the client a mental image of what “a calmer stack” looks like. Additionally, highlight the **Monthly Savings** column by perhaps shading it or using a bold font, since cost savings draw attention.

Design and Branding: Ensure all visuals share a consistent style aligned with syris.' branding – likely clean, minimalistic, with calming colors (perhaps blues/greens). Avoid overly complex graphics that would contradict the “calm” vibe. A few well-chosen charts with flat design and plenty of white space will actually make the report feel high-end. Leading firms like McKinsey often use simplified charts with clear labels and takeaway captions. You might add **brief captions** to each visual in the report that state the insight. For example, below a radar chart: *“Scope focus is above industry norm, indicating a tendency to get granular – meanwhile tempo is below norm, signifying slower pace.”* Such captions ensure the client doesn't misinterpret the graphic and they reinforce the analysis points (this practice is common in consulting reports). According to best practices, **visuals should complement the text, not duplicate it** ³¹ – use them to emphasize or clarify points that would be harder to grasp in text or tables.

In summary, upgrading the calm.profile report with charts and infographics will significantly improve its **readability and impact**. Executives often skim documents, looking for visuals and headings ³² . By providing key insights in a graphical form, you cater to that behavior and increase the perceived value of

the report. The founder will finish reading with not only a list of recommendations, but also a mental picture of their operations (e.g. a red stoplight on “Decision Latency” or a chart of how much time and money they’ll recover). These kinds of visuals make the findings more *tangible*. They also convey a level of professionalism and effort that can set calm.profile apart – for a \$495 diagnostic, the client will feel they got a very robust, data-driven analysis.

5. Refining Diagnostic Questions and Answer Formats (Optional)

While the current 20 questions effectively cover the four workstyle axes, there’s always room to improve clarity and accuracy. Small tweaks in how questions are phrased or how answers are collected can yield more reliable data:

- **Improve Question Clarity:** Ensure each survey item is focused on a single behavior or preference, stated in plain language familiar to agency folks. Avoid double-barreled questions (e.g. “We have clear processes and everyone follows them” – which mixes process clarity and compliance in one item). Instead, split into distinct questions if needed (“We have a clear standard process for projects” vs. “Team members adhere to our defined process”). This makes it easier for respondents to give an accurate answer, thereby increasing reliability ¹². Likewise, frame questions in the context of actual work situations to aid understanding. For example, rather than an abstract statement like “We prioritize structure over flexibility,” ask “*When starting a project, do you prefer to follow a set plan or figure it out as you go?*” – this kind of situational question is easier to relate to, and the answer will clearly indicate the structure orientation.
- **Use Consistent, Calibrated Scales:** If not already, use a Likert-type scale (e.g. 1–5 or 1–7) with labeled endpoints for each question. For instance, 1 = “Strongly disagree” and 5 = “Strongly agree” to a statement, or a slider from “Very unstructured” to “Highly structured”. Consistent scales help respondents get into a rhythm and provide more consistent answers (which improves internal consistency). It also allows you to treat all questions quantitatively similar in the scoring algorithm. Providing anchor descriptions (even if just at the ends) reduces ambiguity in how they interpret the scale ³³. Additionally, consider if a **neutral midpoint** is desired or if forcing a slight lean (even-numbered scale) would produce clearer archetypes – this depends on whether “balanced” is a meaningful category or just noise. Given you allow a “Balanced” variant in axes (e.g. MACRO/MICRO/BALANCED), having an option for neutrality is fine, but ensure it’s used intentionally.
- **Introduce Variety in Question Types (if needed):** All 20 questions need not be uniform. For example, to gauge **collaboration mode**, you might include a ranking or forced-choice question: “Rank these in order of your team’s preferred communication style: (A) In-person or live calls, (B) Scheduled written updates, (C) On-demand chat”. A ranking forces a trade-off decision that can illuminate preferences more starkly than rating everything high. Another idea is a hypothetical scenario: “If a project timeline is slipping, are team members more likely to (A) set up a quick group call, or (B) send updates via email and wait for responses?” – choosing A or B clearly maps to sync vs async tendency. These formats, used sparingly, can enrich the data. However, balance this with user experience; too many complex question types could reduce completion rates. One or two well-designed ones can add depth without overwhelming the respondent.
- **Validate and Iterate Questions:** Since accuracy is key, build in a feedback loop. After a respondent completes the survey (especially in early uses), consider asking a quick follow-up: “Did any question

feel unclear or hard to answer?” Their feedback can pinpoint items to reword. Over time, analyze if any question’s responses don’t correlate well with others on the same axis (as mentioned above). It might indicate that question isn’t a good measure of the intended trait. Replacing it or rephrasing it could increase the reliability of the axis. Also, track the “confidence band” the algorithm produces – if many profiles come out with “low confidence” (meaning the top two archetype scores are close), the survey might need more discriminatory power on certain dimensions. You could address that by adding one or two extra questions targeting the dimensions that often tie, or by widening the scoring difference through weight adjustments.

- **Increase Respondent Engagement:** To improve honest and thoughtful responses, make the assessment as engaging as possible. Keep it concise (20 questions is already reasonably short), and consider adding a brief intro that assures there are no “right or wrong” answers – it’s about understanding their style. Sometimes respondents try to second-guess what is the “best” answer; reassure them that candid answers make the recommendations more useful. The question phrasing can also be framed positively so people don’t feel defensive. For instance, instead of “Our team often misses deadlines (Yes/No)”, phrase it neutrally: “How would you describe your team’s approach to deadlines?” with options from “Highly structured, never miss” to “Flexible, often adjusted”. This avoids respondent bias or reluctance to admit issues, yielding more accurate data. Maintaining a neutral tone in questions, much like in report writing, yields more unbiased results ³⁴ ³⁵ .
- **Optional Team Perspective:** If feasible, involve multiple team members in the diagnostic to increase accuracy. Even if calm.profile is mainly founder-facing, you might allow the founder to forward a link to their senior team to take a subset of the questions (or the whole survey). That way, you get a more 360-degree view. Differences in answers could be insightful (e.g. leadership thinks collaboration is great, but staff feel differently). This could be outside the base scope (and potentially an upsell opportunity), but even a simplified “team pulse” could validate the findings. For example, include a couple of consensus questions: “Rate your agreement: *Our project processes are efficient.*” If the founder rates this high but team average is low, that’s a discussion point in the debrief. It increases the diagnostic’s credibility and depth. Of course, this requires more coordination, so consider it an optional enhancement if automation allows.

In summary, the core 20-question instrument is a solid foundation, but refining question wording and format can boost its accuracy and respondent ease. Think of it as **calibrating an exam**: we want each question to reliably reflect the trait and be easily understood in one way by all respondents. By doing periodic reviews (analyzing responses, soliciting feedback) and tweaking accordingly, the assessment will only get stronger over time. This ultimately feeds back into the report’s quality – the more accurately the questions pinpoint the client’s workstyle and issues, the more on-target your recommendations will be. And a founder who sees themselves accurately “in that mirror” will trust the subsequent prescription and be more inclined to engage further with syris.

By implementing these improvements in report structure, scoring methodology, automation, data visualization, and questionnaire design, **calm.profile** will become a sharper, more efficient diagnostic tool. The deliverable will read like a high-end consulting report – with a clear storyline, strong visual insights, and actionable takeaways – all delivered quickly thanks to automation. Most importantly, it will speak the founder’s language, focusing on their pain points and ROI in a credible, evidence-backed way. This increases the perceived value of the \$495 diagnostic, priming clients for the next steps (calm.stack,

calm.sys) with confidence in your systematic, “calm” approach. The end result: a smoother internal process for you and a more impactful experience for the client, turning the calm.profile from a one-off report into a compelling entry point for transformational work. 36 37

1 2 3 4 5 6 7 9 10 28 30 31 **How to Write and Present a Consulting Report**

<https://slidemodel.com/consulting-report-how-to-write-and-present-one/>

8 36 37 **Revised syris. Brand Positioning_ Embracing the Calm Series.pdf**

<file:///file-GJcXNQZaXYT6eixbbxScfV>

11 32 34 35 **How to Write a Great Business Consulting Report: Best Practices and Report Examples | Databox**

<https://databox.com/business-consulting-report>

12 **[PDF] The Qualtrics handbook of question design - Student Affairs**

<https://studentaffairs.wvu.edu/files/2020-07/ESS%20Assessment%20-%20Survey%20Item%20Guidance.pdf>

13 **Validating a Questionnaire - Sage Research Methods Community**

<https://researchmethodscommunity.sagepub.com/blog/validating-a-questionnaire>

14 15 16 17 20 21 26 27 **dynamic-report-templates.py**

<file:///file-TqbTPtJKxfqeRNfChixCCh>

18 19 22 23 24 25 **Generating PDF Reports: an automation story | by Holistic AI Engineering | Medium**

https://medium.com/@engineering_holistic_ai/generating-pdf-reports-an-automation-story-e8e2a5557893

29 **calm-profile-diagnostic-report.md**

<file:///file-EEYziv92Dc39m94WBILdLW>

33 **How to improve survey data reliability by labeling scales - Kantar**

<https://www.kantar.com/north-america/inspiration/research-services/how-to-improve-survey-data-reliability-by-labeling-scales>